

Ex 1: Glass.data

Cho dữ liệu glass.data.txt

Sử dụng thuật toán ADABoosting/XGBoost & thuật toán cơ sở để dự đoán loại kính dựa trên các thông tin được cung cấp

1. Đọc dữ liệu và gán cho biến data. Xem thông tin data: shape, type, head(), tail(), info. Tiền xử lý dữ liệu (nếu cần)
2. Tạo inputs data với các cột trừ cột type of class, và outputs data với 1 cột là type of class
3. Từ inputs data và outputs data => Tạo X_train, X_test, y_train, y_test với tỷ lệ 70-30
4. Thực hiện ADABoosting/XGBoost với X_train, y_train
5. Dự đoán y từ X_test => so sánh với y_test
6. Đánh giá mô hình => Nhận xét
7. Ghi mô hình (nếu mô hình tốt sau khi đánh giá)

Attribute Information:

1. Id number: 1 to 214
2. RI: refractive index
3. Na: Sodium (unit measurement: weight percent in corresponding oxide, as are attributes 4-10)
4. Mg: Magnesium
5. Al: Aluminum
6. Si: Silicon
7. K: Potassium
8. Ca: Calcium
9. Ba: Barium
10. Fe: Iron
11. Type of glass: (class attribute)

-- 1 building_windows_float_processed -- 2 building_windows_non_float_processed -- 3 vehicle_windows_float_processed -- 4 vehicle_windows_non_float_processed (none in this database) -- 5 containers -- 6 tableware -- 7 headlamps

```
In [1]: # from google.colab import drive
# drive.mount("/content/gdrive", force_remount=True)

In [2]: # %cd '/content/gdrive/My Drive/LDS6_MachineLearning/practice_2023/Chapter8_Boosting/'

In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

In [4]: # import some data to play with
data = pd.read_csv("glass.data.txt", sep=",", header=None)
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 214 entries, 0 to 213
Data columns (total 11 columns):
0      214 non-null int64
1      214 non-null float64
2      214 non-null float64
3      214 non-null float64
4      214 non-null float64
5      214 non-null float64
6      214 non-null float64
7      214 non-null float64
8      214 non-null float64
9      214 non-null float64
10     214 non-null int64
dtypes: float64(9), int64(2)
memory usage: 18.5 KB
```

```
In [5]: data.shape
```

```
Out[5]: (214, 11)
```

```
In [6]: #data.head()
```

```
In [7]: # thống kê số lượng các lớp
data.groupby(10).count()[0]
```

```
Out[7]: 10
1      70
2      76
3      17
5      13
6       9
7      29
Name: 0, dtype: int64
```



```
In [8]: # The columns that we will be making predictions with.
inputs = data.iloc[:,1:-1]
inputs.shape
```

Out[8]: (214, 9)

```
In [9]: inputs.head()
```

Out[9]:

| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|---|---------|-------|------|------|-------|------|------|-----|-----|
| 0 | 1.52101 | 13.64 | 4.49 | 1.10 | 71.78 | 0.06 | 8.75 | 0.0 | 0.0 |
| 1 | 1.51761 | 13.89 | 3.60 | 1.36 | 72.73 | 0.48 | 7.83 | 0.0 | 0.0 |
| 2 | 1.51618 | 13.53 | 3.55 | 1.54 | 72.99 | 0.39 | 7.78 | 0.0 | 0.0 |
| 3 | 1.51766 | 13.21 | 3.69 | 1.29 | 72.61 | 0.57 | 8.22 | 0.0 | 0.0 |
| 4 | 1.51742 | 13.27 | 3.62 | 1.24 | 73.08 | 0.55 | 8.07 | 0.0 | 0.0 |

```
In [10]: # The column that we want to predict.
outputs = data[10]
outputs = np.array(outputs)
outputs.shape
```

Out[10]: (214,)

```
In [11]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(inputs, outputs,
                                                    test_size=0.30,
                                                    random_state=1)
```

Chúng ta không áp dụng AdaBoostClassifier với KNN vì KNeighborsClassifier không hỗ trợ sample_weight (mà trong AdaBoostClassifier cần)

AdaBoost

```
In [12]: from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
# mặc định là DecisionTreeClassifier() nên có thể không cần ghi
ml = DecisionTreeClassifier()
boosting = AdaBoostClassifier(n_estimators=100,
                              base_estimator=ml,
                              learning_rate=1)
```

```
In [13]: # Train model
model_new = boosting.fit(X_train, y_train)
```

```
In [14]: model_new.score(X_train, y_train)
```

Out[14]: 1.0

```
In [15]: model_new.score(X_test, y_test)
```

Out[15]: 0.7384615384615385

```
In [16]: # Kết Luận: Overfitting
```

```
In [17]: from sklearn.ensemble import RandomForestClassifier
ml_1 = RandomForestClassifier(n_estimators=100)
boosting_1 = AdaBoostClassifier(n_estimators=100,
                                base_estimator=ml_1,
                                learning_rate=0.1)
```

```
In [18]: # Train model
boosting_1.fit(X_train, y_train)
```

Out[18]: AdaBoostClassifier(algorithm='SAMME.R',
 base_estimator=RandomForestClassifier(bootstrap=True,
 class_weight=None,
 criterion='gini',
 max_depth=None,
 max_features='auto',
 max_leaf_nodes=None,
 min_impurity_decrease=0.0,
 min_impurity_split=None,
 min_samples_leaf=1,
 min_samples_split=2,
 min_weight_fraction_leaf=0.0,
 n_estimators=100,
 n_jobs=None,
 oob_score=False,
 random_state=None,
 verbose=0,
 warm_start=False),
 learning_rate=0.1, n_estimators=100, random_state=None)


```
In [19]: boosting_1.score(X_train, y_train)
```

```
Out[19]: 1.0
```

```
In [20]: boosting_1.score(X_test, y_test)
```

```
Out[20]: 0.8153846153846154
```

```
In [32]: from sklearn.model_selection import cross_val_score
scores1 = cross_val_score(boosting_1, inputs, outputs, cv=20)
scores1
```

c:\program files\python36\lib\site-packages\sklearn\model_selection_split.py:657: Warning: The least populated class in y has only 9 members, which is too few. The minimum number of members in any class cannot be less than n_splits=20.
% (min_groups, self.n_splits)), Warning)

```
Out[32]: array([0.69230769, 0.69230769, 0.92307692, 0.92307692, 0.76923077,
0.84615385, 0.69230769, 0.76923077, 0.46153846, 0.63636364,
0.9, 0.8, 1., 0.88888889, 0.33333333,
0.66666667, 0.75, 1., 1., 1.])
```

```
In [33]: display(np.mean(scores1), np.std(scores1))
```

```
0.7872241647241648
```

```
0.17651315097478404
```

```
In [23]: # Kết Luận: Vẫn overfitting nhưng có cải thiện hơn
# Còn model nào tốt hơn không? Cho kết quả.
# Thử áp dụng bài toán này với XGBoost.
```

XGBoost

```
In [24]: import xgboost as xgb
```

```
In [25]: xgb_model = xgb.XGBClassifier(random_state=42)
xgb_model.fit(X_train, y_train)
```

```
Out[25]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, gamma=0,
learning_rate=0.1, max_delta_step=0, max_depth=3,
min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
nthread=None, objective='multi:softprob', random_state=42,
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
silent=None, subsample=1, verbosity=1)
```

```
In [26]: xgb_model.score(X_train, y_train)
```

```
Out[26]: 1.0
```

```
In [27]: xgb_model.score(X_test, y_test)
```

```
Out[27]: 0.8307692307692308
```

```
In [30]: from sklearn.model_selection import cross_val_score
scores2 = cross_val_score(xgb_model, inputs, outputs, cv=20)
scores2
```

c:\program files\python36\lib\site-packages\sklearn\model_selection_split.py:657: Warning: The least populated class in y has only 9 members, which is too few. The minimum number of members in any class cannot be less than n_splits=20.
% (min_groups, self.n_splits)), Warning)

```
Out[30]: array([0.69230769, 0.69230769, 0.92307692, 0.76923077, 0.84615385,
0.76923077, 0.53846154, 0.84615385, 0.61538462, 0.72727273,
1., 0.8, 0.8, 0.77777778, 0.11111111,
0.77777778, 0.75, 1., 1., 1.])
```

```
In [31]: display(np.mean(scores2), np.std(scores2))
```

```
0.7718123543123543
```

```
0.19677858024774542
```

```
In [ ]:
```