

## Ex 3: Candy Production

Cho dữ liệu sản xuất kẹo từ 1972 đến 2017 trong tập tin `candy_production.csv`

- Thực hiện việc dự báo sản xuất sản phẩm sử dụng thuật toán ARIMA
- Cho biết trong 12 tháng sau những năm trên thì giá trị bán sản phẩm như thế nào?

```
In [1]: # from google.colab import drive
# drive.mount("/content/gdrive", force_remount=True)
```

```
In [2]: # %cd '/content/gdrive/My Drive/LDS6_MachineLearning/practice/Chapter18_ARIMA/'
```

```
In [3]: import pandas as pd
```

```
In [4]: data = pd.read_csv("candy_production.csv", index_col=0)
data.head()
```

```
Out[4]:
```

IPG3113N	
observation_date	
1972-01-01	85.6945
1972-02-01	71.8200
1972-03-01	66.0229
1972-04-01	64.5645
1972-05-01	65.0100

```
In [5]: data.index = pd.to_datetime(data.index)
```

```
In [6]: data.index.freq = 'MS'
```

```
In [7]: data.index
```

```
Out[7]: DatetimeIndex(['1972-01-01', '1972-02-01', '1972-03-01', '1972-04-01',
                        '1972-05-01', '1972-06-01', '1972-07-01', '1972-08-01',
                        '1972-09-01', '1972-10-01',
                        ...,
                        '2016-11-01', '2016-12-01', '2017-01-01', '2017-02-01',
                        '2017-03-01', '2017-04-01', '2017-05-01', '2017-06-01',
                        '2017-07-01', '2017-08-01'],
                        dtype='datetime64[ns]', name='observation_date', length=548, freq='MS')
```

```
In [8]: data.info()

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 548 entries, 1972-01-01 to 2017-08-01
Freq: MS
Data columns (total 1 columns):
IPG3113N    548 non-null float64
dtypes: float64(1)
memory usage: 8.6 KB
```

```
In [9]: data.columns = ['candy_production']
```

```
In [10]: data.head()
```

```
Out[10]:
```


candy_production	
observation_date	
1972-01-01	85.6945
1972-02-01	71.8200
1972-03-01	66.0229
1972-04-01	64.5645
1972-05-01	65.0100

```
In [11]: import matplotlib.pyplot as plt
```



```
In [12]: data.plot(figsize=(14,8))
```

```
Out[12]: <AxesSubplot:xlabel='observation_date'>
```



```
In [13]: type(data)
```


```
Out[13]: pandas.core.frame.DataFrame
```

```
In [14]: from statsmodels.tsa.seasonal import seasonal_decompose
result = seasonal_decompose(x = data, model='multiplicative')
result
```

```
Out[14]: <statsmodels.tsa.seasonal.DecomposeResult at 0x1bbddc2550>
```

```
In [15]: result.plot()
plt.show()
```

c:\program files\python36\lib\site-packages\pandas\plotting\\_matplotlib\converter.py:245: MatplotlibDeprecationWarning:  
The epoch2num function was deprecated in Matplotlib 3.3 and will be removed two minor releases later.  
base = dates.epoch2num(dt.asi8 / 1.0e9)



```
In [16]: #! pip install pmdarima
```

```
In [17]: from pmdarima import auto_arima
```



```
In [18]: stepwise_model = auto_arima(data, start_p=2, start_q= 2,
                                     max_p=5, max_q=5, m=12,
                                     start_P=1, seasonal=True,
                                     d=1, D=1, trace=True,
                                     error_action='ignore',
                                     suppress_warnings=True,
                                     stepwise=True)
```

```
Performing stepwise search to minimize aic
ARIMA(2,1,2)(1,1,1)[12]      : AIC=2955.432, Time=3.13 sec
ARIMA(0,1,0)(0,1,0)[12]      : AIC=3177.330, Time=0.02 sec
ARIMA(1,1,0)(1,1,0)[12]      : AIC=3067.080, Time=0.20 sec
ARIMA(0,1,1)(0,1,1)[12]      : AIC=2969.534, Time=0.40 sec
ARIMA(2,1,2)(0,1,1)[12]      : AIC=2960.282, Time=1.89 sec
ARIMA(2,1,2)(1,1,0)[12]      : AIC=inf, Time=2.25 sec
ARIMA(2,1,2)(2,1,1)[12]      : AIC=2956.153, Time=10.27 sec
ARIMA(2,1,2)(1,1,2)[12]      : AIC=2956.787, Time=9.57 sec
ARIMA(2,1,2)(0,1,0)[12]      : AIC=inf, Time=0.66 sec
ARIMA(2,1,2)(0,1,2)[12]      : AIC=2954.852, Time=7.14 sec
ARIMA(1,1,2)(0,1,2)[12]      : AIC=2951.313, Time=4.51 sec
ARIMA(1,1,2)(0,1,1)[12]      : AIC=2957.418, Time=1.60 sec
ARIMA(1,1,2)(1,1,2)[12]      : AIC=2953.311, Time=6.35 sec
ARIMA(1,1,2)(1,1,1)[12]      : AIC=2951.707, Time=2.22 sec
ARIMA(0,1,2)(0,1,2)[12]      : AIC=2958.148, Time=2.05 sec
ARIMA(1,1,1)(0,1,2)[12]      : AIC=2954.325, Time=2.67 sec
ARIMA(1,1,3)(0,1,2)[12]      : AIC=2951.129, Time=6.71 sec
ARIMA(1,1,3)(0,1,1)[12]      : AIC=inf, Time=3.05 sec
ARIMA(1,1,3)(1,1,2)[12]      : AIC=2953.126, Time=11.51 sec
ARIMA(1,1,3)(1,1,1)[12]      : AIC=2951.415, Time=3.11 sec
ARIMA(0,1,3)(0,1,2)[12]      : AIC=2960.055, Time=2.61 sec
ARIMA(2,1,3)(0,1,2)[12]      : AIC=2949.061, Time=7.01 sec
ARIMA(2,1,3)(0,1,1)[12]      : AIC=inf, Time=3.19 sec
ARIMA(2,1,3)(1,1,2)[12]      : AIC=2951.061, Time=11.66 sec
ARIMA(2,1,3)(1,1,1)[12]      : AIC=2949.457, Time=4.21 sec
ARIMA(3,1,3)(0,1,2)[12]      : AIC=inf, Time=13.21 sec
ARIMA(2,1,4)(0,1,2)[12]      : AIC=2952.228, Time=12.31 sec
ARIMA(1,1,4)(0,1,2)[12]      : AIC=2951.543, Time=7.16 sec
ARIMA(3,1,2)(0,1,2)[12]      : AIC=2949.561, Time=8.71 sec
ARIMA(3,1,4)(0,1,2)[12]      : AIC=inf, Time=16.53 sec
ARIMA(2,1,3)(0,1,2)[12] intercept : AIC=inf, Time=13.17 sec
```

```
Best model: ARIMA(2,1,3)(0,1,2)[12]
Total fit time: 179.067 seconds
```

```
In [19]: print(stepwise_model.aic())
```

```
2949.061078238775
```

```
In [20]: train = data.loc['1972-01-01':'2009-12-01']
test = data.loc['2009-12-01':]
```

```
In [21]: test
```

```
Out[21]: candy_production
```

observation_date	
2009-12-01	116.5435
2010-01-01	100.3797
2010-02-01	99.0155
2010-03-01	91.9654
2010-04-01	89.4914
...	...
2017-04-01	107.4288
2017-05-01	101.9209
2017-06-01	104.2022
2017-07-01	102.5861
2017-08-01	114.0613

```
93 rows × 1 columns
```

```
In [22]: len(test)
```

```
Out[22]: 93
```

```
In [23]: len(train)
```

```
Out[23]: 456
```

```
In [24]: stepwise_model.fit(train)
```

```
Out[24]: ARIMA(order=(2, 1, 3), scoring_args={}, seasonal_order=(0, 1, 2, 12),
               suppress_warnings=True, with_intercept=False)
```



```
In [25]: future_forecast = stepwise_model.predict(n_periods=len(test))
```

```
In [26]: future_forecast
```

```
Out[26]: array([102.13211266, 100.13213126, 95.28694666, 88.79930746,
 89.64604485, 88.6238308 , 88.81275702, 96.51020015,
105.27723551, 121.83806119, 119.38733708, 118.26943919,
105.1182907 , 101.85950122, 97.03431377, 90.68034887,
90.81467504, 90.608315 , 89.27341719, 97.35147057,
105.69478874, 120.06515434, 119.5472295 , 116.9252781 ,
104.50967307, 101.16107329, 95.95276167, 90.29731377,
89.68479953, 89.99870909, 88.55201676, 96.3075675 ,
105.27373708, 118.95974706, 118.93515196, 116.18425982,
103.49948883, 100.70505237, 94.87074005, 89.68155903,
88.92713193, 89.01866278, 88.06388778, 95.24780982,
104.65334129, 118.18805017, 117.98198654, 115.66670227,
102.46085171, 100.07926131, 94.08731991, 88.75231714,
88.38264899, 87.99999587, 87.43212818, 94.45469015,
103.74534231, 117.61897487, 116.98214476, 115.02855631,
101.65980279, 99.19007972, 93.49582186, 87.77017075,
87.73783175, 87.19256096, 86.55957316, 93.8427827 ,
102.77978515, 116.96731424, 116.16966415, 114.17065258,
101.02935046, 98.24003579, 92.83724093, 86.95380365,
86.89280161, 86.54528703, 85.62400177, 93.17728414,
101.96052959, 116.13356082, 115.50715829, 113.24855232,
100.35700293, 97.4187466 , 92.01333256, 86.27752955,
85.98321321, 85.86621315, 84.80140694, 92.3619398 ,
101.27183248])
```

```
In [27]: future_forecast = pd.DataFrame(future_forecast,
                                     index = test.index,
                                     columns=['Prediction'])
```

```
In [28]: from sklearn.metrics import mean_absolute_error
```

```
In [29]: mean_absolute_error(test['candy_production'], future_forecast)
```

```
Out[29]: 8.524769812375583
```

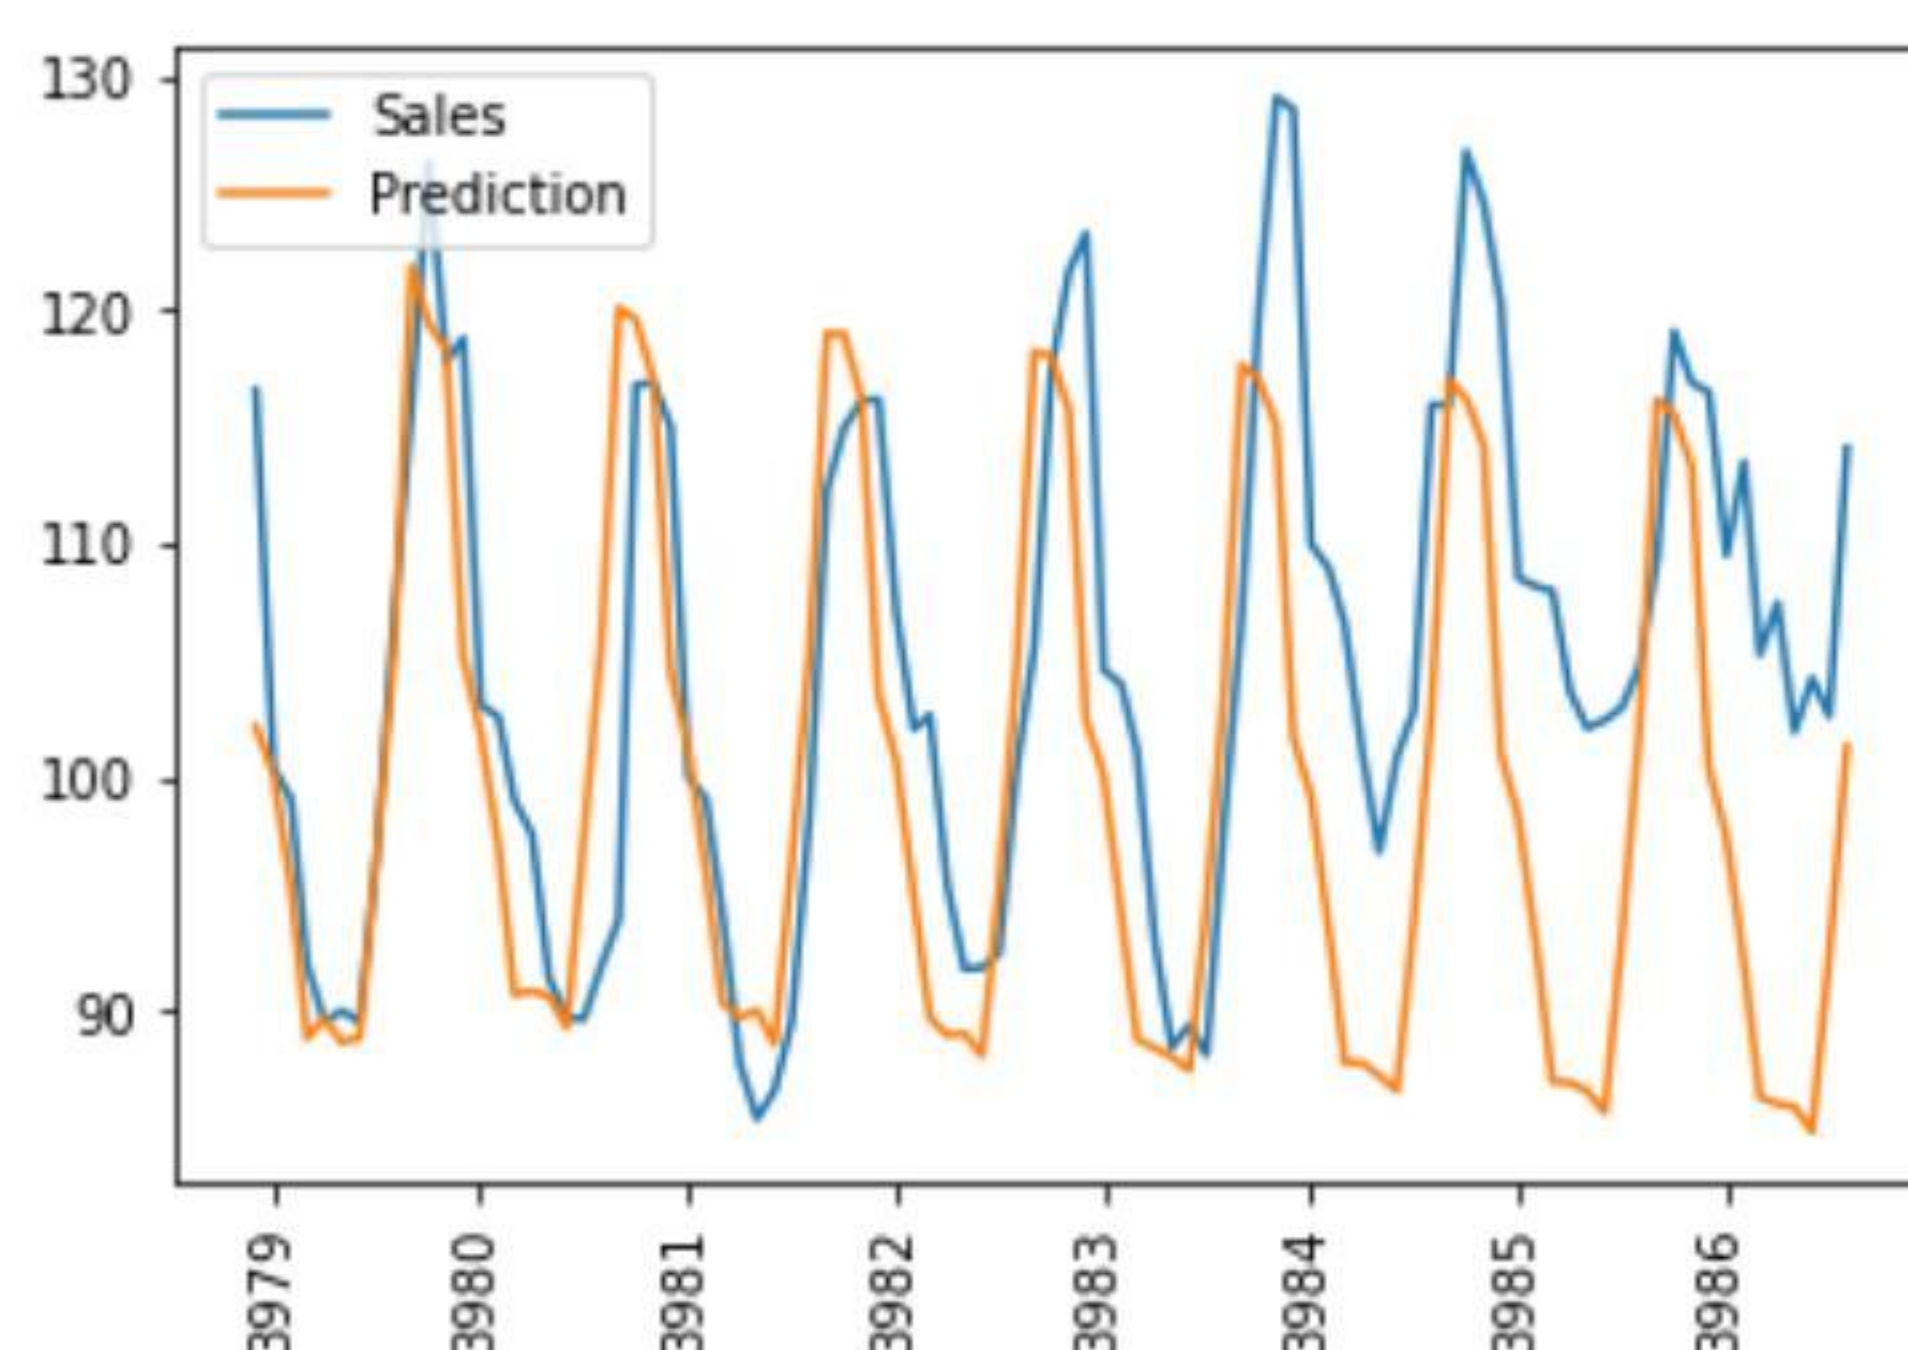
```
In [30]: future_forecast = pd.DataFrame(future_forecast, index = test.index, columns=['Prediction'])
df_merge = test.join(future_forecast)
df_merge.head()
```

```
Out[30]:
```

	candy_production	Prediction
observation_date		
2009-12-01	116.5435	102.132113
2010-01-01	100.3797	100.132131
2010-02-01	99.0155	95.286947
2010-03-01	91.9654	88.799307
2010-04-01	89.4914	89.646045


```
In [31]: plt.plot(test, label='Sales')
plt.plot(future_forecast, label='Prediction')
plt.xticks(rotation='vertical')
plt.legend()
plt.show()
```

c:\program files\python36\lib\site-packages\pandas\plotting\\_matplotlib\converter.py:245: MatplotlibDeprecationWarning:  
The epoch2num function was deprecated in Matplotlib 3.3 and will be removed two minor releases later.  
base = dates.epoch2num(dt.asi8 / 1.0e9)





```
In [32]: plt.figure(figsize=(15,8))
plt.plot(data, label='Acutal')
plt.plot(future_forecast, label='Prediction', color='red')
plt.xticks(rotation='vertical')
plt.legend()
plt.show()
```



```
In [33]: # Dự đoán 12 tháng sau
```

```
In [34]: future_forecast = stepwise_model.predict(n_periods=len(test)+12)
future_forecast
```


```
Out[34]: array([102.13211266, 100.13213126, 95.28694666, 88.79930746,
89.64604485, 88.6238308 , 88.81275702, 96.51020015,
105.27723551, 121.83806119, 119.38733708, 118.26943919,
105.1182907 , 101.85950122, 97.03431377, 90.68034887,
90.81467504, 90.608315 , 89.27341719, 97.35147057,
105.69478874, 120.06515434, 119.5472295 , 116.9252781 ,
104.50967307, 101.16107329, 95.95276167, 90.29731377,
89.68479953, 89.99870909, 88.55201676, 96.3075675 ,
105.27373708, 118.95974706, 118.93515196, 116.18425982,
103.49948883, 100.70505237, 94.87074005, 89.68155903,
88.92713193, 89.01866278, 88.06388778, 95.24780982,
104.65334129, 118.18805017, 117.98198654, 115.66670227,
102.46085171, 100.07926131, 94.08731991, 88.75231714,
88.38264899, 87.99999587, 87.43212818, 94.45469015,
103.74534231, 117.61897487, 116.98214476, 115.02855631,
101.65980279, 99.19007972, 93.49582186, 87.77017075,
87.73783175, 87.19256096, 86.55957316, 93.8427827 ,
102.77978515, 116.96731424, 116.16966415, 114.17065258,
101.02935046, 98.24003579, 92.83724093, 86.95380365,
86.89280161, 86.54528703, 85.62400177, 93.17728414,
101.96052959, 116.13356082, 115.50715829, 113.24855232,
100.35700293, 97.4187466 , 92.01333256, 86.27752955,
85.98321321, 85.86621315, 84.80140694, 92.3619398 ,
101.27183248, 115.23556918, 114.82152429, 112.42526771,
99.54907873, 96.71886061, 91.12606792, 85.58553617,
85.15975584, 85.06468986, 84.09146229, 91.48457837,
100.57370712])
```

```
In [39]: future_forecast[len(test):]
```

```
Out[39]: array([115.23556918, 114.82152429, 112.42526771, 99.54907873,
96.71886061, 91.12606792, 85.58553617, 85.15975584,
85.06468986, 84.09146229, 91.48457837, 100.57370712])
```



```
In [35]: plt.figure(figsize=(8,5))
plt.plot(future_forecast[len(test):], color='red')
plt.xticks(rotation='vertical')
plt.title("Prediction next 12 months")
plt.show()
```



```
In [36]: import numpy as np
```

```
In [37]: plt.plot(np.arange(len(test)+12), future_forecast, color='green')
plt.plot(np.arange(len(test), len(test)+12, 1), future_forecast[len(test):], color='red')
plt.xticks(rotation='vertical')
plt.title("Prediction next 12 months")
plt.show()
```

