



Chapter 17: Linear Regression

Exercise 1: Baseball

Yêu cầu: Áp dụng Line Regression để dự đoán cân nặng dựa trên chiều cao

Cho dữ liệu baseball_2D.txt. Hãy áp dụng Line Regression để dự đoán cân nặng dựa trên chiều cao

- Đọc dữ liệu và gán cho biến data.
- Xem thông tin data: head(), số dòng, số cột, str, summary
- Chuyển dữ liệu theo công thức: chiều cao (m) = chiều cao (inch) * 0.0254, cân nặng (kg) = cân nặng (pound) * 0.453592.
- Vẽ biểu đồ quan sát mối liên hệ giữa inputs và outputs data (scatter plot)
- Kiểm tra outliers => loại outliers
- Tạo train:test từ dữ liệu data với tỉ lệ 70:30
- Thực hiện Linear Regression với X_train, y_train.
- In summary của model
- Dự đoán y_pred từ test => so sánh với y_test
- Tính Mean Square Error (mse)
- Tính Coefficients, Intercept và Variance score
- Cho chiều cao lần lượt: x <- c(1.775, 1.825, 1.925) => dự đoán cân nặng
- Vẽ hình và xem kết quả

```
In [1]: # predict weight based on height
# open and read csv file
data <- read.csv("baseball.csv")
print(head(data))
print(is.data.frame(data))
print(paste("cols", ncol(data)))
print(paste("rows:", nrow(data)))
```

	Name	Team	Position	Height	Weight	Age	PosCategory
1	Adam_Donachie	BAL	Catcher	74	180	22.99	Catcher
2	Paul_Bako	BAL	Catcher	74	215	34.69	Catcher
3	Ramon_Hernandez	BAL	Catcher	72	210	30.78	Catcher
4	Kevin_Millar	BAL	First_Baseman	72	210	35.43	Infielder
5	Chris_Gomez	BAL	First_Baseman	73	188	35.71	Infielder
6	Brian_Roberts	BAL	Second_Baseman	69	176	29.39	Infielder

```
[1] TRUE
[1] "cols 7"
[1] "rows: 1015"
```




In [2]: summary(data)

Name	Team	Position	Height
Chris_Young	NYM	Relief_Pitcher	Min. :67.00
Tony_Peña	ATL	Starting_Pitcher	1st Qu.:72.00
A.J._Burnett	CHC	Outfielder	Median :74.00
A.J._Murray	DET	Catcher	Mean :73.69
A.J._Pierzynski	OAK	Second_Baseman	3rd Qu.:75.00
Aaron_Boone	WAS	First_Baseman	Max. :83.00
(Other)	(Other)	(Other)	

Weight	Age	PosCategory
Min. :150.0	Min. :20.90	Catcher : 76
1st Qu.:186.0	1st Qu.:25.41	Infielder :210
Median :200.0	Median :27.90	Outfielder:194
Mean :201.3	Mean :28.71	Pitcher :535
3rd Qu.:215.0	3rd Qu.:31.19	
Max. :290.0	Max. :48.52	

In [3]: str(data)

```
'data.frame': 1015 obs. of 7 variables:
 $ Name      : Factor w/ 1013 levels "A.J._Burnett",...: 13 778 801 615 199 134
717 703 66 22 ...
 $ Team      : Factor w/ 30 levels "ANA","ARZ","ATL",...: 4 4 4 4 4 4 4 4 4 4
...
 $ Position   : Factor w/ 8 levels "Catcher","First_Baseman",...: 1 1 1 2 2 5 6
8 8 3 ...
 $ Height    : int  74 74 72 72 73 69 69 71 76 71 ...
 $ Weight    : int  180 215 210 210 188 176 209 200 231 180 ...
 $ Age       : num  23 34.7 30.8 35.4 35.7 ...
 $ PosCategory: Factor w/ 4 levels "Catcher","Infielder",...: 1 1 1 2 2 2 2 2 2
3 ...
```

In [4]: `baseball <- data[c("Height", "Weight")]`
`print(head(baseball))`

	Height	Weight
1	74	180
2	74	215
3	72	210
4	72	210
5	73	188
6	69	176




```
In [5]: baseball["Height"] <- baseball["Height"] * 0.0254
baseball["Weight"] <- baseball["Weight"] * 0.453592

print("After preprocessing data:")
print(head(baseball))
```


```
[1] "After preprocessing data:"
      Height  Weight
1 1.8796 81.64656
2 1.8796 97.52228
3 1.8288 95.25432
4 1.8288 95.25432
5 1.8542 85.27530
6 1.7526 79.83219
```

```
In [6]: scatter.smooth(x=baseball$Height, y=baseball$Weight,
                       main="Height vs Weight")
```





```
In [7]: # BoxPlot to Check for outliers
par(mfrow=c(1, 2)) # divide graph area in 2 columns
boxplot(baseball$Height, main="Height",
        sub=paste("Outlier rows: ",
                  boxplot.stats(baseball$Height)$out))
boxplot(baseball$Weight, main="Weight",
        sub=paste("Outlier rows: ",
                  boxplot.stats(baseball$Weight)$out))
```





```
In [8]: # calculate correlation between Width and Length
print(cor(baseball$Height, baseball$Weight))

wt_outliers <- c(boxplot.stats(baseball$Weight)$out)
print("wt_outliers: ")
print(wt_outliers)

ht_outliers <- c(boxplot.stats(baseball$Height)$out)
print("ht_outliers: ")
print(ht_outliers)
```

```
[1] 0.5315393
[1] "wt_outliers: "
[1] 117.9339 122.4698 131.5417 126.0986 117.9339 117.9339 117.9339
[1] "ht_outliers: "
[1] 2.0574 2.0320 2.0320 2.0320 2.0320 2.0828 2.0320 2.0574 2.0828 2.1082
[11] 1.7018 1.7018
```

```
In [9]: #drop rows have outliers
print(paste("Before drop:", nrow(baseball)))
for (record in wt_outliers){
  baseball <- baseball[baseball$Weight != record,]
}
for (record in ht_outliers)
{
  baseball <- baseball[baseball$Height != record,]
}
print(paste("After drop:", nrow(baseball)))
```

```
[1] "Before drop: 1015"
[1] "After drop: 998"
```

```
In [10]: # Create the training (development) and test (validation) data.
# https://rafalab.github.io/dsbook/probability.html#monte-carlo-simulations-for-
set.seed(42) # setting seed to reproduce results of random sampling
trainingRowIndex <- sample(1:nrow(baseball), 0.7*nrow(baseball))
# print("Selected training row indexes:")
# print(trainingRowIndex)
```

```
In [11]: trainingData <- baseball[trainingRowIndex, ] # training data
testData <- baseball[-trainingRowIndex, ] # test data
```

```
In [12]: print("Rows of training data and test data:")
print(nrow(trainingData))
print(nrow(testData))
```

```
[1] "Rows of training data and test data:"
[1] 698
[1] 300
```

```
In [13]: # Develop the model on the training data and use it to predict the Length on tes
lmMod <- lm(Weight ~ Height, data=trainingData) # build the model
```




```
In [14]: iPred <- predict(lmMod, testData) # predict Length
# mean square error according to model
mse <- mean(lmMod$residuals^2)
print(paste("mse: ", mse))
```

```
[1] "mse: 61.1770098543297"
```

```
In [15]: # mean square error of testData
mse_test = mean((testData$Weight - iPred)^2)
print(paste("mse in test: ", mse_test))
```

```
[1] "mse in test: 59.3539332937392"
```

```
In [16]: summary(lmMod)$r.squared
# => r^2 has low value, this model fits ~ 27% data => not good!
```

```
0.2687499653293
```

```
In [17]: # Review diagnostic measures.
print(summary(lmMod)) # model summary
```

Call:

```
lm(formula = Weight ~ Height, data = trainingData)
```

Residuals:

Min	1Q	Median	3Q	Max
-21.3701	-5.8208	0.4391	5.3014	27.8722

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-67.775	9.937	-6.82	1.98e-11	***
Height	85.006	5.315	15.99	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.833 on 696 degrees of freedom

Multiple R-squared: 0.2687, Adjusted R-squared: 0.2677

F-statistic: 255.8 on 1 and 696 DF, p-value: < 2.2e-16

```
In [18]: # model coefficients
print(coef(lmMod) )
# get beta estimate for height
beta_height <- coef(lmMod)["Height"]
print(paste("slope: ",beta_height))
Intercept <- coef(lmMod)["(Intercept)"]
print(paste("Intercept: ",Intercept))
```

(Intercept)	Height
-67.77469	85.00638

```
[1] "slope: 85.0063807905534"
```

```
[1] "Intercept: -67.7746921487327"
```



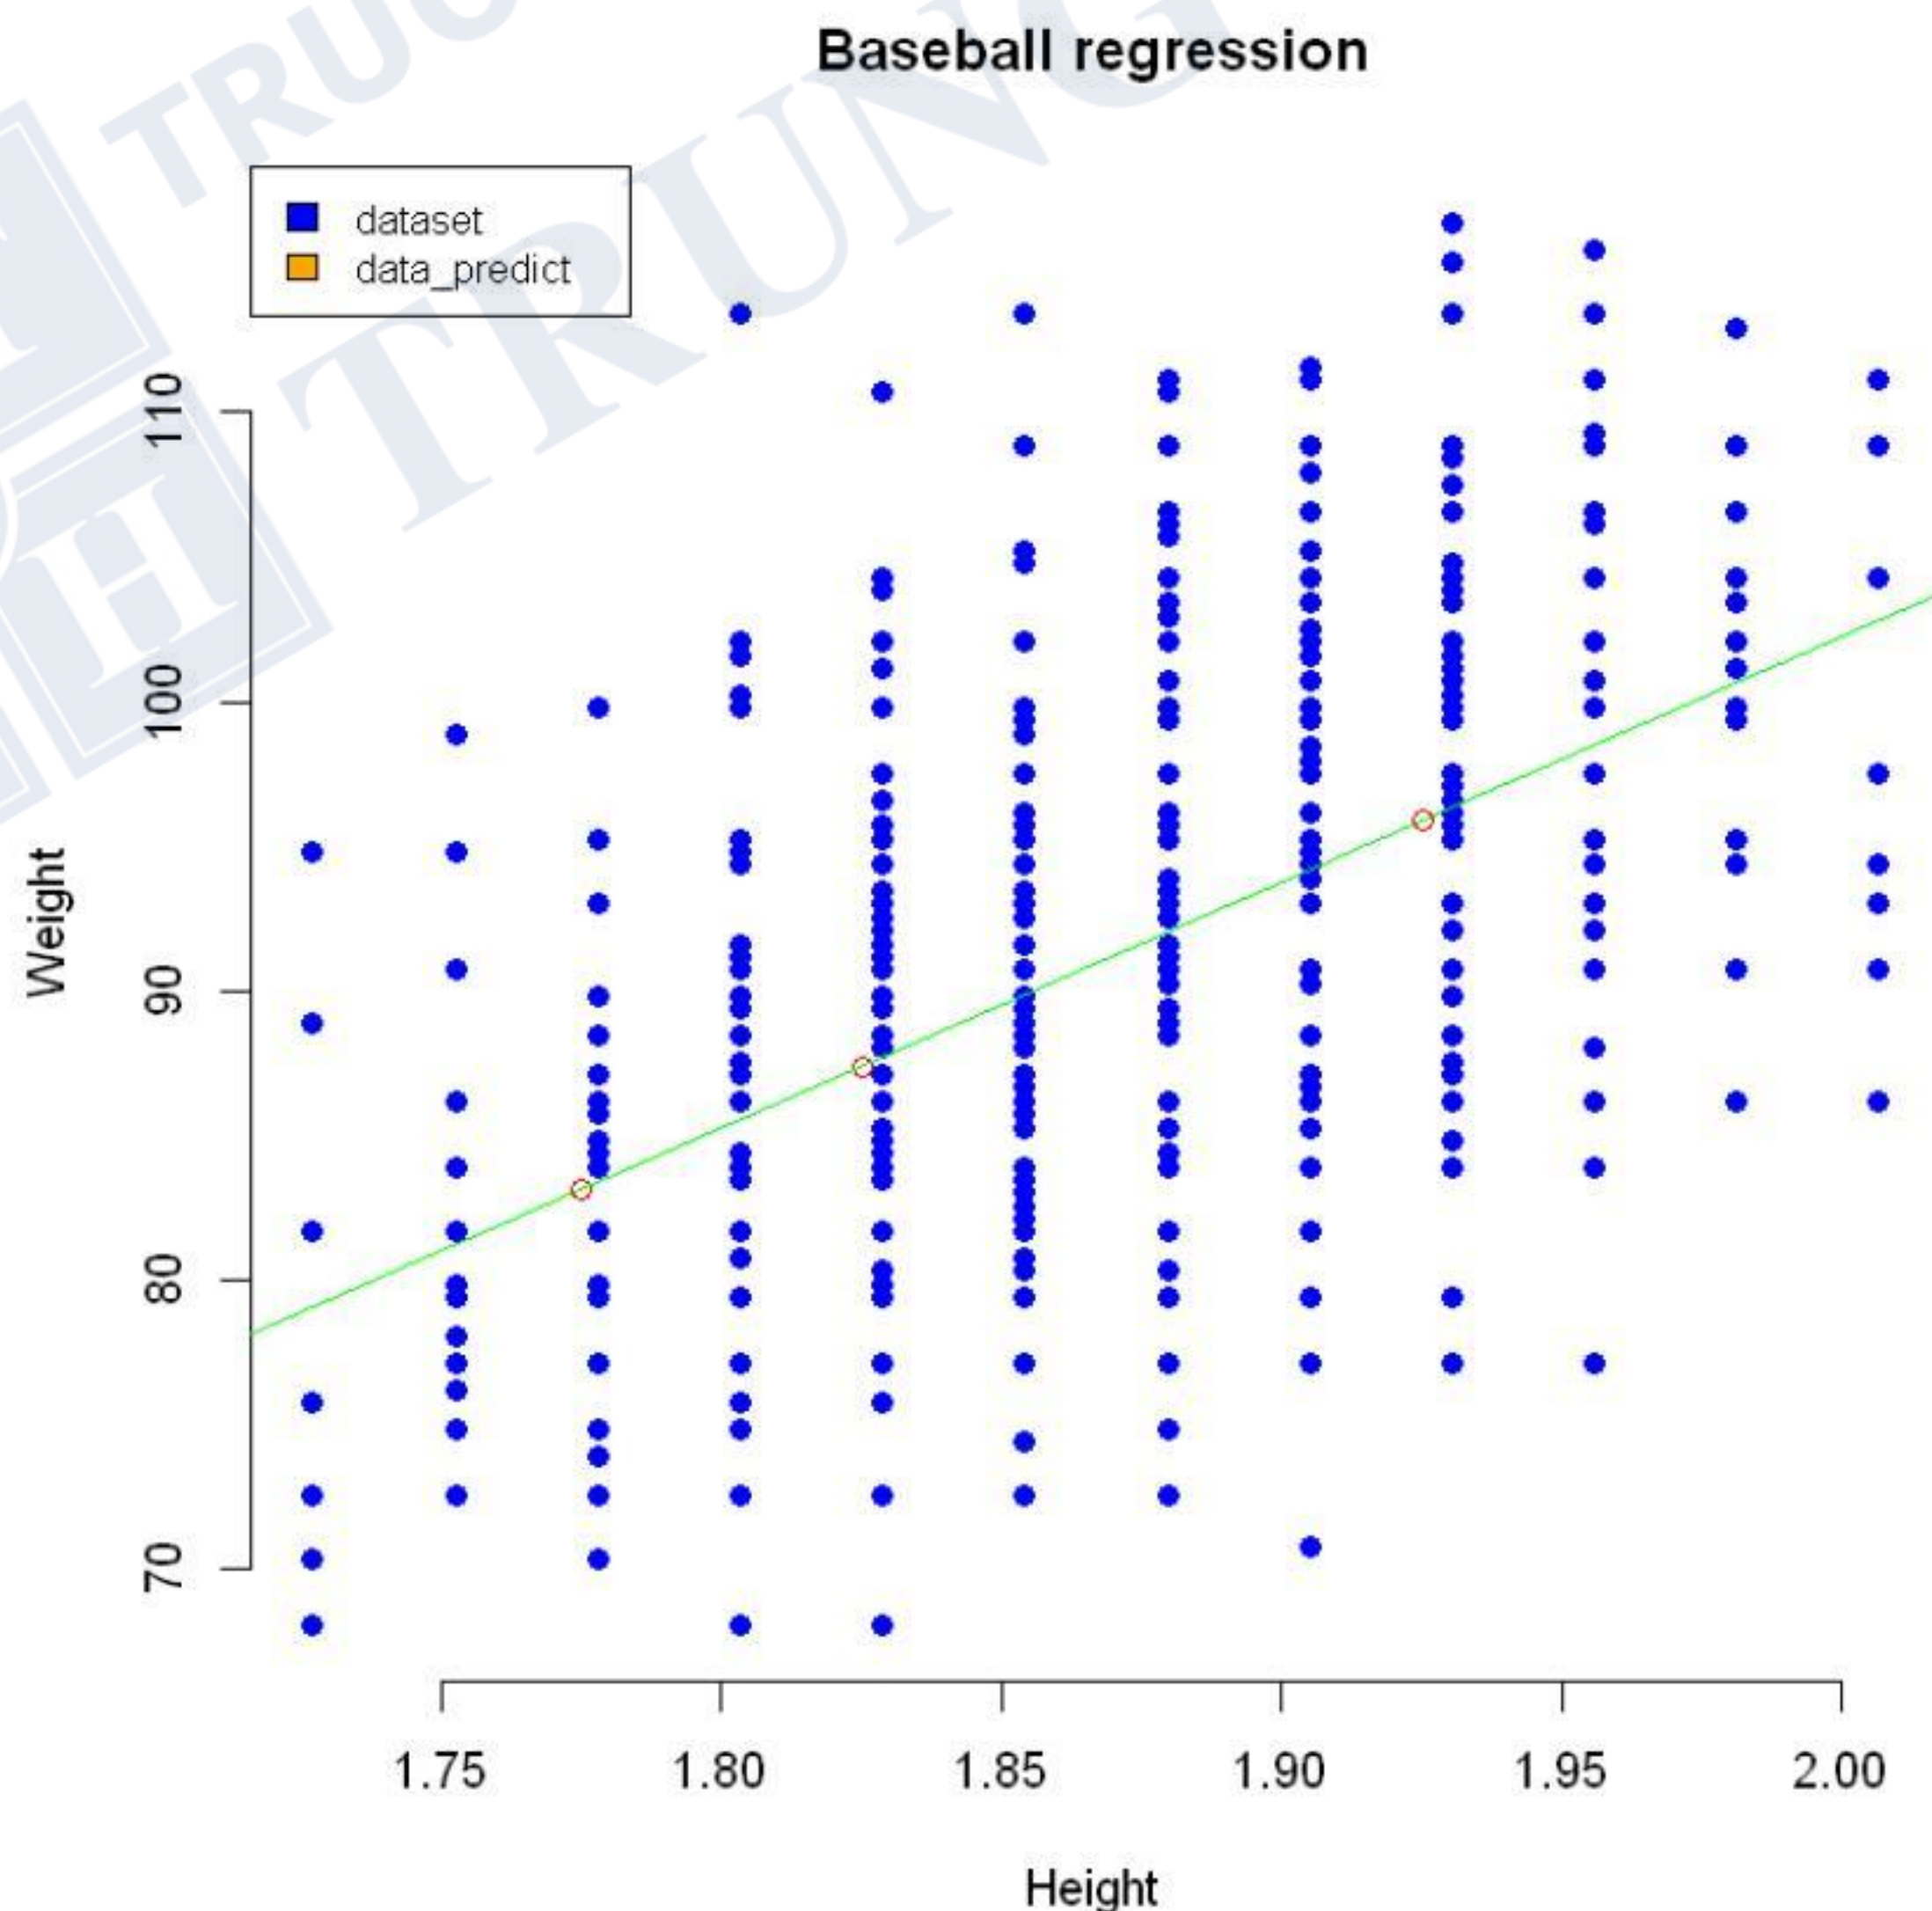

```
In [19]: # new predictions
# solution 1
x <- c(1.775, 1.825, 1.925)
y <- Intercept + beta_height * x
print("Solution 1 - results:")
print(y)
```

```
[1] "Solution 1 - results:"
[1] 83.11163 87.36195 95.86259
```

```
In [20]: # solution 2
y1 <- predict(lmMod, data.frame(Height = x))
print("Solution 2 - results:")
print(y1)
```

```
[1] "Solution 2 - results:"
      1      2      3
83.11163 87.36195 95.86259
```

```
In [21]: # visualization
plot(baseball$Height, baseball$Weight,
     main = "Baseball regression",
     xlab = "Height", ylab = "Weight",
     pch = 19, frame = FALSE, col = 'blue')
lines(x, y, col = 'red', type = 'p')
abline(lmMod, baseball, col = "green")
legend("topleft", c("dataset", "data_predict"),
     cex = 0.8, fill = c("blue", "orange"))
```





TRƯỜNG ĐẠI HỌC KHOA HỌC TỰ NHIÊN
TRUNG TÂM TIN HỌC