



Chapter 8: Tree Methods

Ex1: College

We will be using a college dataset (College.csv) to try to classify colleges as Private or Public based off these features:

Private: A factor with levels No and Yes indicating private or public university
Apps: Number of applications received
Accept: Number of applications accepted
Enroll: Number of new students enrolled
Top10perc: Pct. new students from top 10% of H.S. class
Top25perc: Pct. new students from top 25% of H.S. class
F.Undergrad: Number of fulltime undergraduates
P.Undergrad: Number of parttime undergraduates
Outstate: Out-of-state tuition
Room.Board: Room and board costs
Books: Estimated book costs
Personal: Estimated personal spending
PhD: Pct. of faculty with Ph.D.'s
Terminal: Pct. of faculty with terminal degree
S.F.Ratio: Student/faculty ratio
perc.alumni: Pct. alumni who donate
Expend: Instructional expenditure per student
Grad.Rate: Graduation rate

```
In [1]: import findspark
findspark.init()
```

```
In [2]: from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('treecode').getOrCreate()
```

```
In [3]: # Load training data
data = spark.read.csv('College.csv',inferSchema=True,header=True)
```

```
In [4]: data.count()
```

```
Out[4]: 777
```



In [5]: `data.printSchema()`

```
root
|-- School: string (nullable = true)
|-- Private: string (nullable = true)
|-- Apps: integer (nullable = true)
|-- Accept: integer (nullable = true)
|-- Enroll: integer (nullable = true)
|-- Top10perc: integer (nullable = true)
|-- Top25perc: integer (nullable = true)
|-- F_Undergrad: integer (nullable = true)
|-- P_Undergrad: integer (nullable = true)
|-- Outstate: integer (nullable = true)
|-- Room_Board: integer (nullable = true)
|-- Books: integer (nullable = true)
|-- Personal: integer (nullable = true)
|-- PhD: integer (nullable = true)
|-- Terminal: integer (nullable = true)
|-- S_F_Ratio: double (nullable = true)
|-- perc_alumni: integer (nullable = true)
|-- Expend: integer (nullable = true)
|-- Grad_Rate: integer (nullable = true)
```

In [6]: `data.head()`

Out[6]: Row(School='Abilene Christian University', Private='Yes', Apps=1660, Accept=1232, Enroll=721, Top10perc=23, Top25perc=52, F_Undergrad=2885, P_Undergrad=537, Outstate=7440, Room_Board=3300, Books=450, Personal=2200, PhD=70, Terminal=78, S_F_Ratio=18.1, perc_alumni=12, Expend=7041, Grad_Rate=60)

Spark Formatting of Data

```
In [7]: # It needs to be in the form of two columns
        # ("label", "features")

        # Import VectorAssembler and Vectors
        from pyspark.ml.linalg import Vectors
        from pyspark.ml.feature import VectorAssembler
```



```
In [8]: data.columns
```

```
Out[8]: ['School',  
         'Private',  
         'Apps',  
         'Accept',  
         'Enroll',  
         'Top10perc',  
         'Top25perc',  
         'F_Undergrad',  
         'P_Undergrad',  
         'Outstate',  
         'Room_Board',  
         'Books',  
         'Personal',  
         'PhD',  
         'Terminal',  
         'S_F_Ratio',  
         'perc_alumni',  
         'Expend',  
         'Grad_Rate']
```

```
In [9]: assembler = VectorAssembler(  
        inputCols=['Apps',  
                   'Accept',  
                   'Enroll',  
                   'Top10perc',  
                   'Top25perc',  
                   'F_Undergrad',  
                   'P_Undergrad',  
                   'Outstate',  
                   'Room_Board',  
                   'Books',  
                   'Personal',  
                   'PhD',  
                   'Terminal',  
                   'S_F_Ratio',  
                   'perc_alumni',  
                   'Expend',  
                   'Grad_Rate'],  
        outputCol="features")
```

```
In [10]: output = assembler.transform(data)
```

Deal with Private column being "yes" or "no"

```
In [11]: from pyspark.ml.feature import StringIndexer
```

```
In [12]: indexer = StringIndexer(inputCol="Private", outputCol="PrivateIndex")  
output_fixed = indexer.fit(output).transform(output)
```

```
In [13]: final_data = output_fixed.select("features", 'PrivateIndex')
```



```
In [14]: train_data, test_data = final_data.randomSplit([0.7, 0.3])
```

The Classifiers

```
In [15]: from pyspark.ml.classification import DecisionTreeClassifier
from pyspark.ml.classification import GBTClassifier, RandomForestClassifier
from pyspark.ml import Pipeline
```

Create all three models:

```
In [16]: # Use mostly defaults to make this comparison "fair"

dtc = DecisionTreeClassifier(labelCol='PrivateIndex', featuresCol='features')
rfc = RandomForestClassifier(labelCol='PrivateIndex', featuresCol='features')
gbt = GBTClassifier(labelCol='PrivateIndex', featuresCol='features')
```

Train all three models:

```
In [17]: # Train the models (its three models, so it might take some time)
dtc_model = dtc.fit(train_data)
rfc_model = rfc.fit(train_data)
gbt_model = gbt.fit(train_data)
```

Model Comparison

Let's compare each of these models!

```
In [18]: dtc_predictions = dtc_model.transform(test_data)
rfc_predictions = rfc_model.transform(test_data)
gbt_predictions = gbt_model.transform(test_data)
```

Evaluation Metrics:

```
In [19]: from pyspark.ml.evaluation import MulticlassClassificationEvaluator
```

```
In [20]: # Select (prediction, true label) and compute test error
acc_evaluator = MulticlassClassificationEvaluator(labelCol="PrivateIndex",
                                                  predictionCol="prediction",
                                                  metricName="accuracy")
```

```
In [21]: dtc_acc = acc_evaluator.evaluate(dtc_predictions)
rfc_acc = acc_evaluator.evaluate(rfc_predictions)
gbt_acc = acc_evaluator.evaluate(gbt_predictions)
```



```
In [22]: print("Results:")
print('-'*80)
print('A single decision tree - accuracy: {0:2.2f}%'.format(dtc_acc*100))
print('-'*80)
print('A random forest ensemble - accuracy: {0:2.2f}%'.format(rfc_acc*100))
print('-'*80)
print('A ensemble using GBT - accuracy: {0:2.2f}%'.format(gbt_acc*100))
```

Results:

```
-----
-
A single decision tree - accuracy: 92.51%
-----
-
A random forest ensemble - accuracy: 95.59%
-----
-
A ensemble using GBT - accuracy: 92.95%
```

Optional Assignment - play around with the parameters of each of these models, can you squeeze some more accuracy out of them? Or is the data the limiting factor?