



Chapter 9: K-Means Clustering

Ex1: 3D points

Dataset: data3D.csv.

Requirement:

- Read dataset
- Pre-process data
- Use K-means clustering algorithm to cluster 3D points in data3D.csv.

```
In [8]: import findspark
findspark.init()
```

```
In [9]: import pyspark
```

```
In [10]: from pyspark.sql import SparkSession
```

```
In [11]: spark = SparkSession.builder.appName('kmeans_3D_point').getOrCreate()
```

```
In [12]: # Loads data.
data = spark.read.csv("data3D.csv", header=True,
                      inferSchema=True)
```

```
In [13]: data.show(3)
```

```
+-----+-----+-----+-----+
|   id|          x|          y|          z|
+-----+-----+-----+-----+
|point0|5.647627534046943|-6.356222340123802|-7.240816026826695|
|point1|4.414367138680041|-10.32624175635328| 8.963324308916228|
|point2|5.005396944639823|-9.301070062115645| 10.35473056351597|
+-----+-----+-----+-----+
only showing top 3 rows
```

```
In [14]: from pyspark.sql.functions import col
```

```
In [17]: data = data.select(['x', 'y', 'z'])
```



```
In [18]: data.show(3)
```

```
+-----+-----+-----+
|          x|          y|          z|
+-----+-----+-----+
|5.647627534046943|-6.356222340123802|-7.240816026826695|
|4.414367138680041|-10.32624175635328| 8.963324308916228|
|5.005396944639823|-9.301070062115645| 10.35473056351597|
+-----+-----+-----+
only showing top 3 rows
```

Format from data

```
In [19]: from pyspark.ml.linalg import Vectors
         from pyspark.ml.feature import VectorAssembler
```

```
In [20]: data.columns
```

```
Out[20]: ['x', 'y', 'z']
```

```
In [21]: vec_assembler = VectorAssembler(inputCols = data.columns,
                                         outputCol='features')
```

```
In [25]: final_data = vec_assembler.transform(data)
```

Scale the Data

```
In [22]: from pyspark.ml.feature import StandardScaler
```

```
In [23]: scaler = StandardScaler(inputCol="features",
                                outputCol="scaledFeatures",
                                withStd=True,
                                withMean=False)
```

```
In [26]: # Compute summary statistics by fitting the StandardScaler
         scalerModel = scaler.fit(final_data)
```

```
In [27]: # Normalize each feature to have unit standard deviation.
         final_data = scalerModel.transform(final_data)
```



In [28]: `final_data.show(3, False)`

```
+-----+-----+-----+-----+
|x          |y          |z          |features
|scaledFeatures|          |          |
+-----+-----+-----+-----+
|5.647627534046943|-6.356222340123802|-7.240816026826695|[5.647627534046943,-6.356222340123802,-7.240816026826695]|[1.0159673512169811,-0.81335799160424,-1.1300631023636807]|
|4.414367138680041|-10.32624175635328|8.963324308916228|[4.414367138680041,-10.32624175635328,8.963324308916228]|[0.7941127247055396,-1.3213715327025133,1.3988923401033817]|
|5.005396944639823|-9.301070062115645|10.35473056351597|[5.005396944639823,-9.301070062115645,10.35473056351597]|[0.9004347126254774,-1.1901880174546189,1.6160469899240197]|
+-----+-----+-----+-----+
only showing top 3 rows
```

Train the Model and Evaluate

Select k with minimum WSSSE: k between 2 - 10

In [29]: `from pyspark.ml.clustering import KMeans`

In [31]: `# Trains a k-means model.`


```
k_list = []
wssse_list = []
for k in range(2,11):
    kmeans = KMeans(featuresCol='scaledFeatures',k=k)
    model = kmeans.fit(final_data)
    wssse = model.computeCost(final_data)
    k_list.append(k)
    wssse_list.append(wssse)
print("With k =", k, "Set Sum of Squared Errors = " + str(wssse))
```

```
With k = 2 Set Sum of Squared Errors = 1155067.2563008904
With k = 3 Set Sum of Squared Errors = 297656.40920043044
With k = 4 Set Sum of Squared Errors = 146718.50451770602
With k = 5 Set Sum of Squared Errors = 72720.18504204818
With k = 6 Set Sum of Squared Errors = 68583.68208541229
With k = 7 Set Sum of Squared Errors = 64483.499921190276
With k = 8 Set Sum of Squared Errors = 60180.384358530224
With k = 9 Set Sum of Squared Errors = 58364.71279295459
With k = 10 Set Sum of Squared Errors = 52251.866281251554
```



```
In [32]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [33]: plt.plot(k_list, wssse_list)
plt.show()
```



- According to Elbow Method, we choose $k = 5$. Look like there is very little gain after $k=5$, so we stick to that choice when processing the full data set.

Select $k = 5$

```
In [34]: # Trains a k-means model.
kmeans = KMeans(featuresCol='scaledFeatures', k=5)
model = kmeans.fit(final_data)
```

```
In [35]: # Evaluate clustering by computing Within Set Sum of Squared Errors.
wssse = model.computeCost(final_data)
print("Within Set Sum of Squared Errors = " + str(wssse))
```

Within Set Sum of Squared Errors = 72720.18504204816

```
In [36]: # Shows the result.
centers = model.clusterCenters()
print("Cluster Centers: ")
for center in centers:
    print(center)
```

```
Cluster Centers:
[-1.59030525  0.93719373  0.31581855]
[ 1.19780483 -0.7365009  -0.99420908]
[ 0.74896823 -1.2269015   1.46702482]
[-0.45182036  1.15367577  0.72369935]
[ 0.35333753 -0.87989287 -1.0735622 ]
```



```
In [37]: predictions = model.transform(final_data)
```

```
In [38]: predictions.select("prediction").show(5)
```

```
+-----+
|prediction|
+-----+
|         1|
|         2|
|         2|
|         2|
|         0|
+-----+
```

only showing top 5 rows

```
In [ ]: # Check number points of each cluster
```

```
In [60]: predictions.groupBy('prediction').count().show()
```

```
+-----+-----+
|prediction| count|
+-----+-----+
|         1|199987|
|         3|200017|
|         4|200013|
|         2|200000|
|         0|199983|
+-----+-----+
```

```
In [ ]: # Our clustering algorithm created 5 equally sized clusters with K=5
```

```
In [39]: data_result = predictions.select("prediction")
         data_result.columns
```

```
Out[39]: ['prediction']
```

```
In [40]: type(data_result)
```

```
Out[40]: pyspark.sql.dataframe.DataFrame
```



In [41]: `final_data.show(3, False)`

```
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+
|x          |y          |z          |features
|scaledFeatures          |
+-----+-----+-----+-----+
+-----+
|5.647627534046943|-6.356222340123802|-7.240816026826695|[5.647627534046943,-6.
356222340123802,-7.240816026826695]|[1.0159673512169811,-0.81335799160424,-1.13
00631023636807]|
|4.414367138680041|-10.32624175635328|8.963324308916228|[4.414367138680041,-1
0.32624175635328,8.963324308916228]|[0.7941127247055396,-1.3213715327025133,1.
3988923401033817]|
|5.005396944639823|-9.301070062115645|10.35473056351597|[5.005396944639823,-9.
301070062115645,10.35473056351597]|[0.9004347126254774,-1.1901880174546189,1.6
160469899240197]|
+-----+-----+-----+-----+
+-----+
only showing top 3 rows
```

In [42]: `temp = final_data.select("scaledFeatures").rdd.map(lambda x: \
x[0].toArray().tolist()).toDF()`

In [43]: `temp.show(3)`

```
+-----+-----+-----+
|          _1|          _2|          _3|
+-----+-----+-----+
|1.0159673512169811| -0.81335799160424|-1.1300631023636807|
|0.7941127247055396|-1.3213715327025133| 1.3988923401033817|
|0.9004347126254774|-1.1901880174546189| 1.6160469899240197|
+-----+-----+-----+
only showing top 3 rows
```

In [44]: `import pyspark.sql.functions as f`

In [45]: `# since there is no common column between these two dataframes add row_index so
temp=temp.withColumn('row_index', f.monotonically_increasing_id())
data_result=data_result.withColumn('row_index',
f.monotonically_increasing_id())
temp = temp.join(data_result,
on=["row_index"]).sort("row_index").drop("row_index")`



In [46]: `temp.show(3)`

```
+-----+-----+-----+-----+
|          _1|          _2|          _3|prediction|
+-----+-----+-----+-----+
|1.0159673512169811| -0.81335799160424|-1.1300631023636807|      1|
|0.7941127247055396|-1.3213715327025133| 1.3988923401033817|      2|
|0.9004347126254774|-1.1901880174546189| 1.6160469899240197|      2|
+-----+-----+-----+-----+
only showing top 3 rows
```

In [48]: `temp = temp.select(col("_1").alias("x_scale"),
col("_2").alias("y_scale"),
col("_3").alias("z_scale"),
"prediction")`

In [49]: `df = temp.toPandas()`

In [54]: `df.head(3)`

Out[54]:

	x_scale	y_scale	z_scale	prediction
0	1.015967	-0.813358	-1.130063	1
1	0.794113	-1.321372	1.398892	2
2	0.900435	-1.190188	1.616047	2

In [50]: `centers_df = pd.DataFrame(centers)
centers_df.head()`

Out[50]:


	0	1	2
0	-1.590305	0.937194	0.315819
1	1.197805	-0.736501	-0.994209
2	0.748968	-1.226902	1.467025
3	-0.451820	1.153676	0.723699
4	0.353338	-0.879893	-1.073562

In [55]: `import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D`



```
In [58]: threedee = plt.figure(figsize=(12,10)).gca(projection='3d')
threedee.scatter(df.x_scale, df.y_scale,
                 df.z_scale,
                 c=df.prediction)
threedee.set_xlabel('x')
threedee.set_ylabel('y')
threedee.set_zlabel('z')

plt.show()
```



Combine results



```
In [51]: # since there is no common column between these two dataframes add row_index
final_data=final_data.withColumn('row_index',
                                f.monotonically_increasing_id())
temp=temp.withColumn('row_index',
                    f.monotonically_increasing_id())
final_data = final_data.join(temp,
                            on=["row_index"]).sort("row_index").drop("row_index")
```

```
In [52]: final_data.show(3, False)
```

```
+-----+-----+-----+-----+
|x          |y          |z          |features          |
|scaledFeatures          |x_scale          |
|y_scale          |z_scale          |prediction|
+-----+-----+-----+-----+
|5.647627534046943|-6.356222340123802|-7.240816026826695|[5.647627534046943,-6.356222340123802,-7.240816026826695]|[1.0159673512169811,-0.81335799160424,-1.1300631023636807]|1.0159673512169811|-0.81335799160424|-1.1300631023636807|1
|4.414367138680041|-10.32624175635328|8.963324308916228|[4.414367138680041,-10.32624175635328,8.963324308916228]|[0.7941127247055396,-1.3213715327025133,1.3988923401033817]|0.7941127247055396|-1.3213715327025133|1.3988923401033817|2
|5.005396944639823|-9.301070062115645|10.35473056351597|[5.005396944639823,-9.301070062115645,10.35473056351597]|[0.9004347126254774,-1.1901880174546189,1.6160469899240197]|0.9004347126254774|-1.1901880174546189|1.6160469899240197|2
only showing top 3 rows
```